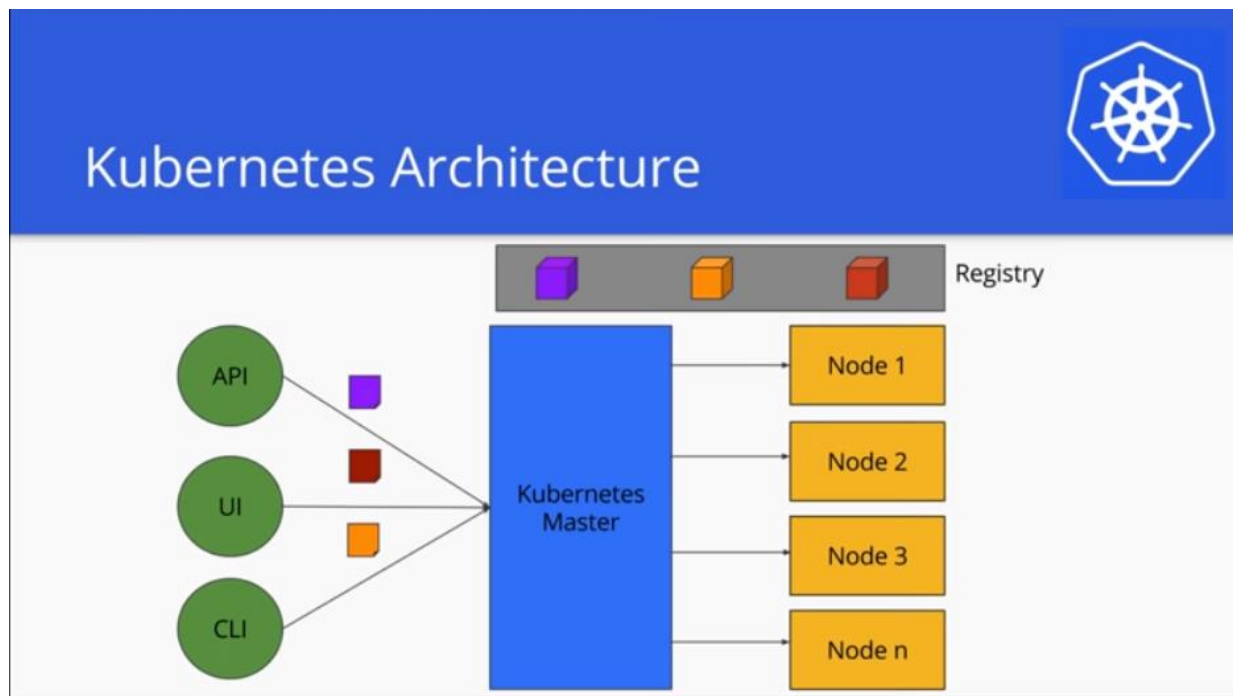


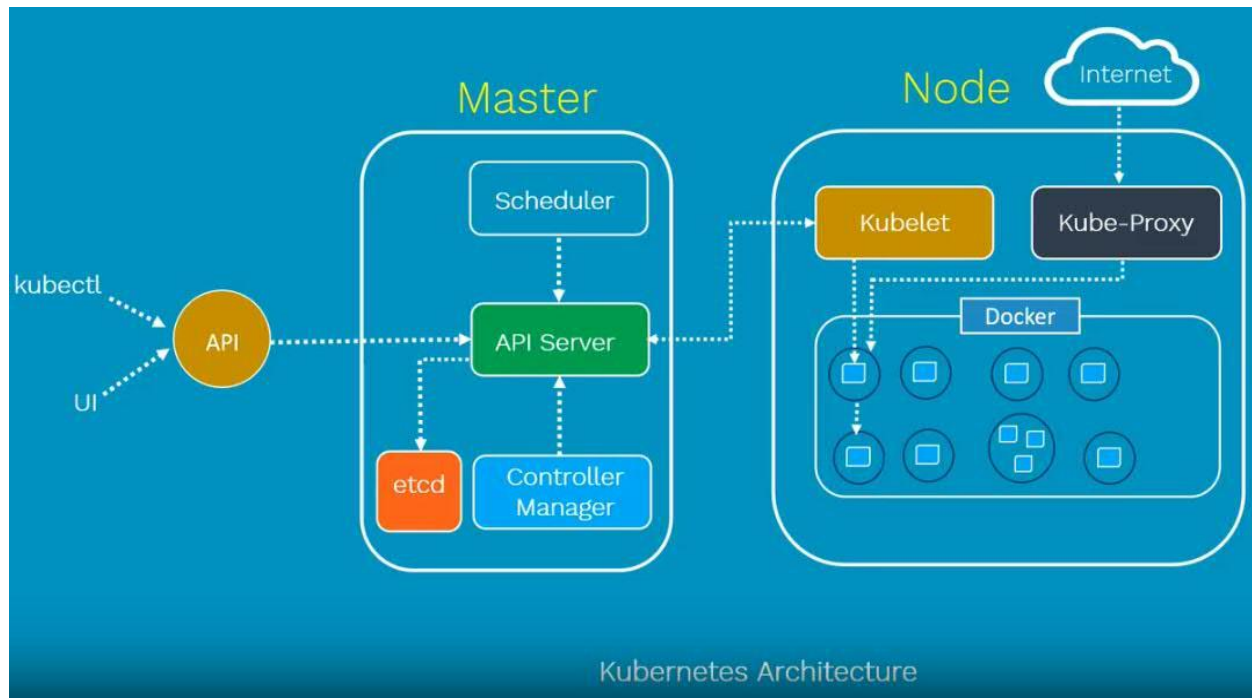
Kubernetes

Features of Kubernetes

Following are some of the important features of Kubernetes.

- Continues development, integration and deployment
- Containerized infrastructure
- Application-centric management
- Auto-scalable infrastructure
- Environment consistency across development testing and production
- Loosely coupled infrastructure, where each component can act as a separate unit
- Higher density of resource utilization
- Predictable infrastructure which is going to be created.





Pod scaling:

Pods Tight coupling and Loose coupling:

Pods life cycle: Pending, Running, Succeeded, Failed:

Pod:

```

apiVersion: v1
kind: Pod
metadata:
  name: hello-pod → name of the pod
spec:
  containers:
    - name: hello-ctrl – name of the container
      image: Jenkins → name of the images
      ports: - container port number
              - containerPort: 8080

```

1. NodePort
2. ClusterIP
3. LoadBalancer

```
# kubectl get nodes
# kubectl get create -f pod.yml
# kubectl get pods
# kubectl describe pods
# kubectl get pods -o wide
# kubectl get pods/hello-pod
# kubectl get pods --all-namespaces
# kubectl delete pods/hello-pod
```


Replication Controller:

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: nginx-rc
spec:
  replicas: 3
  selector:
    app: nginx
  template:
    metadata:
      name: nginx
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
```

```
# kubectl scale rc nginx-rc --replicas=5
```

```
# kubectl get po -l app=nginx-app
```

Service:



apiVersion: v1

kind: Service

metadata:

name: nodeport-svc

labels:

app: nginx-app

spec:

selector:

app: nginx-app

type: NodePort

ports:

- nodePort: 31000

port: 80

targetPort: 80

apiVersion: v1

kind: Service

metadata:

name: nodeport-svc

labels:

app: nginx-app

spec:

selector:

app: nginx-app

type: LoadBalancer

ports:

- nodePort: 31000

port: 80

targetPort: 80

ReplicaSet:

Equality Based and **Set-Based**



```
...  
selector:  
  app: nginx  
  tier: frontend  
...
```

=

```
...  
selector:  
  matchLabels:  
    app: nginx  
    tier: frontend  
...
```

Supports on Older Resources such as:

- ReplicationControllers,
- Services

Supports on newer resources such as:

- ReplicaSets
- Deployments
- Jobs
- DaemonSet

Deployment:

1. Recreate
2. RollingUpdate
3. Canary
4. Blue/Green

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-deploy

labels:

app: nginx-app

spec:

replicas: 3

selector:

```
matchLabels:
  app: nginx-app
template:
  metadata:
    labels:
      app: nginx-app
  spec:
    containers:
      - name: nginx-pod
        image: nginx:1.7.9
        ports:
          - containerPort: 80
```

We can update using in 2 ways. Using **kubectl set** and **kubectl edit** commands.

1st option: using kubectl set command.

kubectl set image deploy nginx-deploy nginx-pod =nginx:1.9.1

Or

kubectl set image deploy nginx-deploy nginx- pod =nginx:1.9.1 --record

2nd option: using kubectl edit command.

kubectl edit deploy nginx-deployment

#kubectl rollout history deployment/nginx-deployment

Scale Up and down:

kubectl scale deployment deploy nginx-deployment --replicas=5

Kubernetes Cluster creation:

Step1: Create required no.of instances. Minimum 2 should require (1 master and 1 slave).

Step2:

Let's disable SELinux and SWAP. To disable those, our instances needs few port numbers to be opened on all the instances. (6443, 10250).

Execute below on all nodes (master and all worker).

```
# systemctl stop firewalld
```

```
# systemctl disable firewalld
```

```
# swapoff -a
```

```
# sed -i.bak -r 's/(.+ swap .+)/#\1/' /etc/fstab
```

```
# setenforce 0
```

```
# sed -i 's/enforcing/disabled/g' /etc/selinux/conf
```

Step3: Now it's time to install required packages for Kubernetes:

```
# cat <<EOF > /etc/yum.repos.d/kubernetes.repo
```

```
[kubernetes]
```

```
name=Kubernetes
```

```
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
```

```
enabled=1
```

```
gpgcheck=1
```

```
repo_gpgcheck=1
```

```
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg
```

```
https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
```

```
exclude=kube*
```

```
EOF
```

```
# yum update -y
```

If the above command doesn't work, then open the file by using below command and paste the content and save, else ignore. Then try the above command. So it will work.

Opening file: `vi /etc/yum.repos.d/kubernetes.repo`

Content to save:

```
[kubernetes]
```

```
name=Kubernetes
```

```
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64/
```

```
enabled=1
```



```
gpgcheck=1
repo_gpgcheck=0
gpgkey=https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
```

yum install -y docker kubeadm kubelet kubect1 --disableexcludes=Kubernetes

If Docker is not installed using above command, then install it manually using below commands:

```
sudo yum install yum-utils
```

```
sudo yum-config-manager --enable rhui-REGION-rhel-server-extras
```

```
sudo yum install docker -y
```

systemctl enable docker && systemctl start docker

systemctl enable kubelet && systemctl start kubelet

cat <<EOF > /etc/sysctl.d/k8s.conf

net.bridge.bridge-nf-call-ip6tables = 1

net.bridge.bridge-nf-call-iptables = 1

EOF

sysctl net.bridge.bridge-nf-call-iptables=1

If above command doesn't work, then first run below command and then try:

modprobe br_netfilter sudo sysctl -p

sysctl net.ipv4.ip_forward=1

sysctl --system

systemctl daemon-reload

systemctl restart kubelet

We can copy and run all the commands at once and run on master node. If it works fine, then run the same commands on all worker nodes.

Step4: Lets configure Kubernetes master node (run below command only on master node):

kubeadm init --pod-network-cidr=10.240.0.0/16

Then copy the generated commands to a note pad:

Now run below commands on master to make kubect1 to work.

Run few(below) commands on master node:

```
# mkdir -p $HOME/.kube
```

```
# sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
# sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

To make the network work between the nodes on Kubernetes configuration, we can use different plugins like weave, flannel and calico.

Lets try to add flannel plugin:

```
# kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/v0.9.1/Documentation/kube-flannel.yml
```

If the above command doesn't work, then try below one.

```
# kubectl apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/k8s-manifests/kube-flannel-rbac.yml
```

If you want to use weave plugin instead of flannel, then use below command.

```
kubectl apply --filename https://git.io/weave-kube-1.6
```

These network plugin commands make's DNS up.

Make sure the DNS is up or not by running below command.

To confirm pod network is successfully installed or not, run below command.

```
# kubectl get pods --all-namespaces
```

Step5: Lets join all worker nodes to cluster by running the join command on all worker nodes.

Repeat the same step on all worker nodes.

Step6: Testing.

To make sure on which worker node our pod is running. User below command.

```
# kubectl get nodes
```

Ingress:

apiVersion: extensions/v1beta1

kind: Ingress

metadata:

name: example-ingress

annotations:

ingress.kubernetes.io/rewrite-target: /

spec:

rules:

- http:

paths:

- path: /jenkins

backend:

serviceName: jenkins-svc

servicePort: 8080

- path: /nginx

backend:

serviceName: nginx-svc

servicePort: 80